

Advanced Programming

Debugging: Debuggers

Example - Debugging in Docker Container in VSCode:

- Click 'Run and Debug' again - the code will stop at the breakpoint, we can inspect the variables in the panel on the left
- Can use the icons at the top to step over/into/out
- If you add a breakpoint, two new windows will appear on the left:
 - Call Stack - shows which functions have been called to reach the line with the breakpoint
 - Watch - allows you to type in variable names that you want to watch throughout the execution

Advanced Programming

Debugging: GDB

- We can also run GDB from the command line (installed via Dockerfile):
 - `cd /usr/src/dockertest1/`
 - `gdb data_class_reader_buggy`
- We then add breakpoints/watchpoints and control the program flow using command line
 - The sample session (<https://sourceware.org/gdb/current/onlinedocs/gdb.html/Sample-Session.html#Sample-Session>) is a useful starting point

Advanced Programming

Debugging: GDB

- Add breakpoint with eg. `break data_reader_class_buggy.cpp:47`
- Run with `run` (or `r`)
- Use `n` for next, `s` to step into a subroutine, `c` to continue, `p` to print a value
- `bt` gives a backtrace to see where we are in the stack

```
(gdb) break data_reader_class_buggy.cpp:47
Breakpoint 1 at 0x26b7: file /usr/src/dockertest1/data_reader_class_buggy.cpp, line 48.
(gdb) run
Starting program: /usr/src/dockertest1/data_reader_class_buggy
warning: Error disabling address space randomization: Operation not permitted
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Breakpoint 1, main () at /usr/src/dockertest1/data_reader_class_buggy.cpp:48
48      CSVReader reader(filename);
(gdb) █
```

Advanced Programming

Debugging: GDB

Other useful features are:

- Calling functions from command line eg. `call calculateAverage()`
- Adding watch points eg. `watch average`
 - Program must be running and stopped at a breakpoint
- Inspecting the code without running eg. look at specific memory address, `1 *0x8000000000000000`
- Inspect core dumps eg. to find out which line caused a seg fault
- Define your own commands via `~/.gdbinit`

Advanced Programming

Debugging: Valgrind

- Valgrind is open source 'instrumentation framework' for building dynamic analysis tools
- This includes profiling tools (see later) and debugging
- For debugging, Valgrind is most useful for detecting memory-related errors and threading errors for parallel applications
- Most popular (and default) Valgrind tool is `memcheck`
- To use Valgrind, first compile your program with `-g` flag to include debugging information
 - This will ensure that error messages include exact line numbers
- Do not compile with any optimisation flags higher than `-O0`

Advanced Programming

Debugging: Valgrind

- *Example: run memory check using*

```
valgrind --leak-check=yes ./data_reader_class_buggy
```

(this enables the detailed memory leak detector)

- To change from default tool `memcheck`, add `--tool=`
- Program is run on a `synthetic' CPU provided by Valgrind core
- `Memcheck` adds code to check every memory access and computed value
- Note: program will run 20-30x slower and use a lot more memory

Advanced Programming

Debugging: Valgrind

- Example output for `data_reader_class_buggy`:

```
● root@526151d9bd1f:/usr/src/dockertest1# valgrind --leak-check=yes ./data_reader_class_buggy
==4655== Memcheck, a memory error detector
==4655== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==4655== Using Valgrind-3.18.1 and LibVEX; rerun with -h for copyright info
==4655== Command: ./data_reader_class_buggy
==4655==
Average value of the second column: -nan
==4655==
==4655== HEAP SUMMARY:
==4655==    in use at exit: 0 bytes in 0 blocks
==4655==   total heap usage: 3 allocs, 3 frees, 74,200 bytes allocated
==4655==
==4655== All heap blocks were freed -- no leaks are possible
==4655==
==4655== For lists of detected and suppressed errors, rerun with: -s
==4655== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
○ root@526151d9bd1f:/usr/src/dockertest1#
```

Advanced Programming

Debugging: Valgrind

- For more detailed debugging, Valgrind can interface with gdb

Example: run the command

```
valgrind --vgdb-error=0 ./data_reader_class_buggy
```

then follow instructions provided in the terminal

- Start gdb in another shell
- Tunnel to vgdb (output from first command is input to second command)
- Can debug with gdb commands and combine with eg. the command `monitor` to use Valgrind tools

Advanced Programming

Profiling

Advanced Programming

Profiling

- Profiling is the measuring of the performance of a program
- Can identify where the program spends its time/memory
- Usually used to identify bottlenecks to enable optimisation
- It can also help to identify bugs, if certain functions are being called more or less than expected
- Profiling can be carried out by:
 - Putting timers manually into the program
 - Using software eg. gprof, Valgrind, VTune
 - Note: there is no straightforward profiling tool built into VSCode, as there is for debugging (other than for JavaScript)

Advanced Programming

Profiling

- Performance is primarily measured in two ways:
 - Wall time - total time elapsed while a certain section of the code is running
 - Processor time
 - You will often see the phrase 'CPU time' used as a measurement
 - Nowadays, GPUs are becoming increasingly popular, so 'GPU time' will likely also be widely used
 - The concept behind both is the same - the total time taken by the processor to run part of the code
 - This excludes communication time, I/O and other costs

Advanced Programming

Profiling

- Simple profiling can be done using the Unix `time` command
 - `real` - total time for program to run from start to finish i.e. wall time
 - `user` - CPU time spent in user mode
 - `sys` - CPU time spent in kernel mode (access hardware components etc.)
 - *Example: Try this on one of the examples from the class repo*

```
root@526151d9bd1f:/usr/src/dockertest1# time ./data_reader_class_buggy
Average value of the second column: -nan

real    0m0.020s
user    0m0.013s
sys     0m0.007s
root@526151d9bd1f:/usr/src/dockertest1#
```

Advanced Programming

Profiling

- For more detailed profiling, you will need to use additional tools:
 - Timers - most coding languages will have a timer module that can be used to write timers into the code
 - Profilers - these are programs that can tell the user where the time running code is being spent, eg. functions, wait time for parallel programs
- (There is a parallel here with print debugging vs software debuggers)
- In both cases, timing data can fluctuate - usually a good idea to take an average/median

Advanced Programming

Profiling: Instrumenting with Timers

Advantages:

- Straightforward to implement
- Straightforward to run and interpret
- You can specify exactly which parts you want to time

Disadvantages:

- Takes time to implement in a large code
- Can easily miss sections if implemented sporadically
- Can provide incomplete information eg. Wall time vs CPU time

Advanced Programming

Profiling: Instrumenting with Timers

- Most up to date timer for C++ is `<chrono>` - measures wall time
- Another commonly used library in C++ and C is `time.h` (`<ctime>`)
 - Can measure CPU (`<time.h>`) or wall time (`<sys/time.h>`)
 - Use of this library is sometimes discouraged (out of date, not thread safe)
- Equivalent timers exist in other languages eg. Python `time` and Fortran `CPU_TIME`

Advanced Programming

Profiling: Instrumenting with Timers

- `<chrono>` - main functions used are `high_resolution_clock` and `duration_cast` (converts time into desired measurement, eg. milliseconds)

Example 1: Compile `chrono_timer.cpp` and run a few times

- `time.h` - main function is `clock()` and macro `CLOCKS_PER_SEC`

Example 2: Compile `timeh_time.cpp` and run a few times

Advanced Programming

Profiling: Profilers

Advantages:

- Can quickly identify bottlenecks, especially useful for parallel code
- Some have user-friendly graphical interfaces

Disadvantages:

- Need to learn how to use
- Sometimes unavailable on large machines, need to arrange installation
- Can add overhead which skews performance data

Advanced Programming

Profiling: gprof

- `gprof` (GNU profiler) is an open source profiler for Unix applications
- It provides the user with (https://ftp.gnu.org/old-gnu/Manuals/gprof-2.9.1/html_mono/gprof.html)
 - *Flat profile* - how much time the program spends in each function, and how many times that function is called
 - *Call graph* - which functions called each function, and which other functions they called. This also estimates how much time is spent in the subroutines of each function
 - *Annotated source listing* - copy of program's source code, labelled with number of times each line is executed, `-A` flag

Advanced Programming

Profiling: gprof

- Before profiling with `gprof`, the program must be compiled with the flag `-pg`
 - The `-g` flag can also be useful for line-by-line profiling and basic-block counting
- The program must be run before profiling to generate the information for `gprof`
 - It will run slower due to the time taken to collect and write the profile
- Running a program compiled with the `-pg` flag will generate a file called `gmon.out` - this is the profiling information
- We can then run `gprof` with e.g. `gprof chrono_timer > output`

Example: Run on eg. chrono_timer.cpp and examine output

Advanced Programming

Profiling: gprof

```
Call graph (explanation follows)

granularity: each sample hit covers 4 byte(s) for 0.33% of 3.02 seconds

index % time    self  children   called    name
-----
[1]   100.0    0.00    3.02             <spontaneous>
      1.76    0.00      1/1      main [1]
      0.88    0.00      1/1      someFunction() [2]
      0.38    0.00      1/1      anotherFunction() [3]
      0.00    0.00      3/3      yetAnotherFunction() [4]
      0.00    0.00      3/3      std::common_type<std::chrono::duration<long, std::ratio<1l, 1000000000l> >, std::chrono::dur\
ation<long, std::ratio<1l, 1000000000l> >>::type std::chrono::operator-(<std::chrono::_V2::system_clock, std::chrono::duration<long, std::rat\
io<1l, 1000000000l> >, std::chrono::duration<long, std::ratio<1l, 1000000000l> > >(std::chrono::time_point<std::chrono::_V2::system_clock, st\
d::chrono::duration<long, std::ratio<1l, 1000000000l> > > const&, std::chrono::time_point<std::chrono::_V2::system_clock, std::chrono::durati\
on<long, std::ratio<1l, 1000000000l> > > const&) [18]
      0.00    0.00      3/3      std::enable_if<std::chrono::__is_duration<std::chrono::duration<long, std::ratio<1l, 1000l> \
> >::value, std::chrono::duration<long, std::ratio<1l, 1000l> >>::type std::chrono::duration_cast<std::chrono::duration<long, std::ratio<1l,\
1000l> >, long, std::ratio<1l, 1000000000l> >(std::chrono::duration<long, std::ratio<1l, 1000000000l> > const&) [14]
      0.00    0.00      3/3      std::chrono::duration<long, std::ratio<1l, 1000l> >::count() const [13]
-----
[2]   58.3    1.76    0.00      1/1      main [1]
      1.76    0.00      1      someFunction() [2]
-----
[3]   29.1    0.88    0.00      1/1      main [1]
      0.88    0.00      1      anotherFunction() [3]
-----
[4]   12.6    0.38    0.00      1/1      main [1]
      0.38    0.00      1      yetAnotherFunction() [4]
-----
      0.00    0.00      3/9      std::chrono::duration<long, std::ratio<1l, 1000l> > std::chrono::__duration_cast_impl<std::c\
hrono::duration<long, std::ratio<1l, 1000l> >, std::ratio<1l, 1000000000l>, long, true, false>::__cast<long, std::ratio<1l, 1000000000l> >(std::\
chrono::duration<long, std::ratio<1l, 1000000000l> > const&) [15]
      0.00    0.00      6/9      std::common_type<std::chrono::duration<long, std::ratio<1l, 1000000000l> >, std::chrono::dur\
ation<long, std::ratio<1l, 1000000000l> >>::type std::chrono::operator-(<long, std::ratio<1l, 1000000000l>, long, std::ratio<1l, 1000000000l>\
>(std::chrono::duration<long, std::ratio<1l, 1000000000l> > const&, std::chrono::duration<long, std::ratio<1l, 1000000000l> > const&) [19]
[11]   0.0    0.00    0.00      9      std::chrono::duration<long, std::ratio<1l, 1000000000l> >::count() const [11]
-----
      0.00    0.00      6/6      std::common_type<std::chrono::duration<long, std::ratio<1l, 1000000000l> >, std::chrono::dur\
ation<long, std::ratio<1l, 1000000000l> >>::type std::chrono::operator-(<std::chrono::_V2::system_clock, std::chrono::duration<long, std::rat\
```

Advanced Programming

Profiling: Valgrind for Profiling

- We can also use Valgrind for profiling
- Run command such as eg.

```
valgrind --tool=callgrind ./chrono_timer
```

(Remember we need to have compiled using `-g` to use valgrind)

- Note that with no optimisation, this is significantly slower than gprof
- The output depends significantly on the level of optimisation - we use `-O0` here
- Outputs `callgrind.out.xxxxx` which can be interpreted using `callgrind_annotate callgrind.out.xxxxx`

Example: try out the above commands

Advanced Programming

Profiling: Valgrind for Profiling

```
-- Auto-annotated source: chrono_timer.cpp
-----
Ir
.
.      #include <iostream>
.      #include <chrono>
.
.      // Function to be timed
.      void someFunction()
3 ( 0.00%) {
.          // Simulating some work
3,000,000,004 (58.80%) for (int i = 0; i < 1000000000; ++i)
.          {
.              // Do some computation
.          }
4 ( 0.00%) }
.
.      // Another function to be timed
.      void anotherFunction()
3 ( 0.00%) {
.          // Simulating some work
1,500,000,004 (29.40%) for (int i = 0; i < 500000000; ++i)
.          {
.              // Do some computation
.          }
4 ( 0.00%) }
.
.      // Yet another function to be timed
.      void yetAnotherFunction()
3 ( 0.00%) {
.          // Simulating some work
600,000,004 (11.76%) for (int i = 0; i < 200000000; ++i)
.          {
.              // Do some computation
.          }
4 ( 0.00%) }
.
.      int main()
```

Advanced Programming

Profiling: Other Advanced Tools

- Advanced profiling tools usually require licenses, eg. Intel VTune
- Provides much more detailed information than open source options
- Particularly useful for parallel code eg. effectiveness of threading and vectorisation, scalability
- Often will have a GUI - these can be tricky to set up remotely, but your system admin should be able to help

Advanced Programming

Profiling: Other Advanced Tools

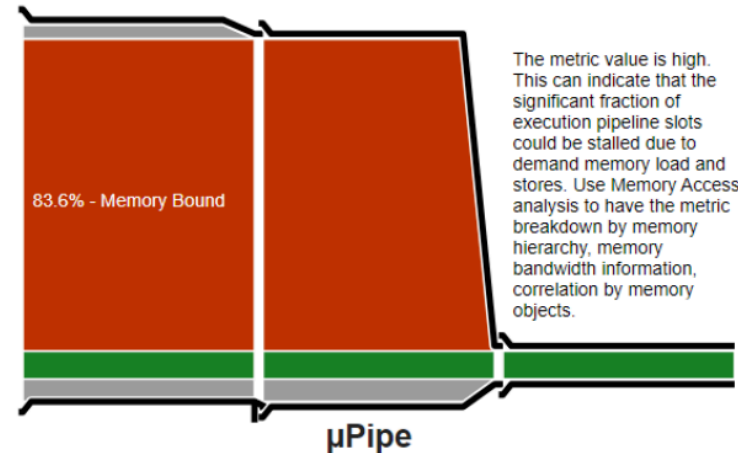
Microarchitecture Exploration Microarchitecture Exploration ? 🔍

Analysis Configuration Collection Log **Summary** Bottom-up Event Count Platform

INTEL VTUNE PROFILER

Elapsed Time [Ⓢ]: 2.731s

Clockticks:	37,686,500,000
Instructions Retired:	3,773,400,000
CPI Rate [Ⓢ] :	9.987 🔻
MUX Reliability [Ⓢ] :	0.923
Retiring [Ⓢ] :	7.5% of Pipeline Slots
Front-End Bound [Ⓢ] :	3.7% of Pipeline Slots
Bad Speculation [Ⓢ] :	0.0% of Pipeline Slots
Back-End Bound [Ⓢ] :	89.6% 🔻 of Pipeline Slots
Memory Bound [Ⓢ] :	83.6% 🔻 of Pipeline Slots
L1 Bound [Ⓢ] :	0.0% of Clockticks
L2 Bound [Ⓢ] :	1.0% of Clockticks
L3 Bound [Ⓢ] :	19.1% 🔻 of Clockticks
DRAM Bound [Ⓢ] :	66.2% 🔻 of Clockticks
Memory Bandwidth [Ⓢ] :	81.8% 🔻 of Clockticks
Memory Latency [Ⓢ] :	11.5% 🔻 of Clockticks
Store Bound [Ⓢ] :	0.0% of Clockticks
Core Bound [Ⓢ] :	6.0% of Pipeline Slots
Average CPU Frequency [Ⓢ] :	2.7 GHz
Total Thread Count:	11
Paused Time [Ⓢ] :	0s



This diagram represents inefficiencies in CPU usage. Treat it as a pipe with an output flow equal to the "pipe efficiency" ratio: (Actual Instructions Retired)/(Maximum Possible Instruction Retired). If there are pipeline stalls decreasing the pipe efficiency, the pipe shape gets more narrow.

Effective Physical Core Utilization [Ⓢ]: 63.4% (2.538 out of 4) 🔻

Effective Logical Core Utilization [Ⓢ]: 63.4% (5.076 out of 8) 🔻

Effective CPU Utilization Histogram

This histogram displays a percentage of the wall time the specific number of CPUs were running simultaneously. Spin and Overhead time adds to the Idle CPU utilization value.

